

STATS 305A_HW 4

Muhammad Ahmed Chaudhry

12/9/2021

In the knitted markdown, some code lines get cropped out due to page margins. However, the R markdown and knitted files may also be viewed at <https://github.com/Ahmed14974/STATS305A>

Problem 4.1 c)

We run the `generate-outlier-data.r` file separately which generates our data for us with a train-test split and also provides us with a function to minimize the objective function in equation (4.1).

```
set.seed(2022)
data <- generate.data()
```

```
# Getting the training and test matrices
X_train <- as.matrix(data[[1]])
y_train <- as.matrix(data[[2]])
X_test  <- as.matrix(data[[3]])
y_test  <- as.matrix(data[[4]])
```

Next, to implement the leave one out cross validation procedure, we use the `minimize.robust.ridge` function provided and build from it:

```
iters <- seq(-2, 1, length.out = 25)

lambdas <- 10^iters

l_lambda <- c()

for (j in 1:length(lambdas)){
  beta_hat <- minimize.robust.ridge(X_train, y_train, lambda = lambdas[j])

  epsilon_hat <- y_train - (X_train %*% beta_hat)

  l_prime <- ((exp(epsilon_hat)) / (1+exp(epsilon_hat))) - ((exp(-1 * epsilon_hat)) / ((1+exp(-1 * epsilon_hat))))
  l_2prime <- ((exp(epsilon_hat)) / ((1+exp(epsilon_hat))^2)) + ((exp(-1 * epsilon_hat)) / ((1+exp(-1 * epsilon_hat))^2))

  XTL <- (t(X_train)) %*% (diag(as.vector(l_2prime)))

  H <- ((1 / (dim(X_train)[1])) * (XTL %*% X_train) + (lambdas[j] * diag(dim(X_train)[2])))

  H_inv <- solve(H)

  # H_k <- list()
  H_k_inv <- list()
  g_k <- list()
}
```

```

delta_hat <- list()
epsilon_hat_neg_k <- c()
loss <- c()

for (i in 1:dim(X_train)[1]){
  numerator <- ((l_2prime[i]) / dim(X_train)[1]) * (H_inv %%% X_train[i,]) %%% (t(X_train[i,]) %%% H_inv)
  denominator <- as.numeric(1 - ((l_2prime[i]/dim(X_train)[1]) * (t(X_train[i,]) %%% H_inv) %%% (X_train[i,]) %%% t(X_train[i,])))
  # H_k_iter <- H - ( (1 / dim(X_train)[1]) * (l_2prime[i]) * (X_train[i,] %%% t(X_train[i,])) )
  # H_k[[i]] <- H_k_iter
  H_k_inv_iter <- H_inv + (numerator/denominator)
  H_k_inv[[i]] <- H_k_inv_iter
  g_k_iter <- (l_prime[i]/(dim(X_train)[1])) * (X_train[i,])
  g_k[[i]] <- g_k_iter
  delta_hat[[i]] <- (-1) * (H_k_inv_iter %%% g_k_iter)
  eps_hat_k <- epsilon_hat[i] + as.numeric((-1) * (X_train[i,] %%% delta_hat[[i]]))
  epsilon_hat_neg_k <- append(epsilon_hat_neg_k, eps_hat_k)
  loss_iter <- log(1 + exp(eps_hat_k)) + log(1 + exp(-1 * eps_hat_k))
  loss <- append(loss, loss_iter)
}

l_lambda <- append(l_lambda, mean(loss))
}

plot_data <- as.data.frame(cbind(iters, l_lambda))


plot_1 <- ggplot(plot_data,
  aes(x = iters, y = l_lambda)) +
  geom_point(aes(x = iters, y = l_lambda))+
  geom_line(aes(x = iters,
    y = l_lambda))+
  xlab("log(lambda)") +
  ylab("L00 error") +

  theme_bw() +
  theme(axis.text.x=element_text(size=8),
    panel.grid.major = element_blank(),
    panel.grid.minor = element_blank())+
  ggtitle("Plot 1. L00 error against log(lambda)")
  # geom_text(
  #   label=round(plot_data$avg_gap,5),
  #   vjust = -3.0,
  #   # hjust = 1.0,
  #   position = position_dodge(width=1),
  #   size =2
  # )

plot_1

```

Plot 1. LOO error against log(lambda)



From the plot above we see that the leave-one-out error first decreases as $\log(\lambda)$ increases from -2 to just above -1, and then starts to increase again as $\log(\lambda)$ increases to 1.

d) From the analysis above, we get the $\hat{\lambda}$ that minimizes the LOO error, from which we can get ERR_{test}

```
lambda_hat <- 10^iters[which.min(l_lambda)]

beta_hat_new <- minimize.robust.ridge(X_train, y_train, lambda = lambda_hat)

error <- y_test - (X_test %*% beta_hat_new)

ERR_test <- median(abs(error))
```

Now, we fit the least squares regularized regression:

```
iters <- seq(log10(2), log10(50), length.out = 25)

lambdas <- 10^iters

ERR_test_ls_lambda <- c()

for (i in 1:length(lambdas)){
  XTX <- t(X_train) %*% X_train

  XTXlambda <- XTX + (lambdas[i] * diag(dim(X_train)[2]))

  XTXlambda_inv <- solve(XTXlambda)

  beta_hat_ls <- (XTXlambda_inv) %*% (t(X_train) %*% y_train)
```

```

error_ls <- y_test - (X_test %*% beta_hat_ls)

ERR_test_ls <- median(abs(error_ls))

ERR_test_ls_lambda <- append(ERR_test_ls_lambda, ERR_test_ls)
}

ERR_diff <- ERR_test_ls_lambda - ERR_test

plot_data_2 <- as.data.frame(cbind(lambdas, ERR_diff))


plot_2 <- ggplot(plot_data_2,
  aes(x = lambdas, y = ERR_diff)) +
  geom_point(aes(x = lambdas, y = ERR_diff))+
  geom_line(aes(x = lambdas, y = ERR_diff))+
  xlab("log(lambda)") +
  ylab("ERR_test_ls - ERR_test") +

  theme_bw() +
  theme(axis.text.x=element_text(size=8),
    panel.grid.major = element_blank(),
    panel.grid.minor = element_blank())+
  ggtitle("Plot 2. Error difference against log(lambda), 10 outliers")
  # geom_text(
  #   label=round(plot_data$avg_gap,5),
  #   vjust = -3.0,
  #   # hjust = 1.0,
  #   position = position_dodge(width=1),
  #   size =2
  # )

plot_2

```

Plot 2. Error difference against $\log(\lambda)$, 10 outliers



We see that as the regularization increases, the error difference decreases towards 0. This is because with greater regularization, the parameters are penalized enough to not be affected by outliers and so the median absolute prediction error gets closer and closer to that from the tuned robust ridge regression.

e) Now, we repeat the same experiment as in part d) but vary the number of outliers, and generate the plots needed:

0 outliers:

```
# rm(list=ls())
# source('generate-outlier-data.r')

# Generating the data
set.seed(2022)
data <- generate.data(num.outliers = 0)

# Getting the training and test matrices
X_train <- as.matrix(data[[1]])
y_train <- as.matrix(data[[2]])
X_test <- as.matrix(data[[3]])
y_test <- as.matrix(data[[4]])

# Tuning the robust ridge regression
iters <- seq(-2, 1, length.out = 25)

lambdas <- 10^iters
```

```

l_lambda <- c()

for (j in 1:length(lambdas)){
  beta_hat <- minimize.robust.ridge(X_train, y_train, lambda = lambdas[j])

  epsilon_hat <- y_train - (X_train %>% beta_hat)

  l_prime <- ((exp(epsilon_hat)) / (1+exp(epsilon_hat))) - ((exp(-1 * epsilon_hat)) / ((1+exp(-1 * epsilon_hat)))
  l_2prime <- ((exp(epsilon_hat)) / ((1+exp(epsilon_hat))^2)) + ((exp(-1 * epsilon_hat)) / ((1+exp(-1 * epsilon_hat))^2))

  XTL <- (t(X_train)) %>% (diag(as.vector(l_2prime)))

  H <- ((1) / (dim(X_train)[1])) * (XTL %>% X_train) + (lambdas[j] * diag(dim(X_train)[2]))

  H_inv <- solve(H)

  # H_k <- list()
  H_k_inv <- list()
  g_k <- list()
  delta_hat <- list()
  epsilon_hat_neg_k <- c()
  loss <- c()

  for (i in 1:dim(X_train)[1]){
    numerator <- ((l_2prime[i]) / dim(X_train)[1]) * (H_inv %>% X_train[i,]) %>% (t(X_train[i,]) %>% H_inv)
    denominator <- as.numeric(1 - ((l_2prime[i]/dim(X_train)[1]) * (t(X_train[i,]) %>% H_inv) %>% (X_train[i,]) %>% t(X_train[i,])))
    # H_k_iter <- H - ( (1 / dim(X_train)[1]) * (l_2prime[i]) * (X_train[i,] %>% t(X_train[i,])) )
    # H_k[[i]] <- H_k_iter
    H_k_inv_iter <- H_inv + (numerator/denominator)
    H_k_inv[[i]] <- H_k_inv_iter
    g_k_iter <- (l_prime[i]/(dim(X_train)[1])) * (X_train[i,])
    g_k[[i]] <- g_k_iter
    delta_hat[[i]] <- (-1) * (H_k_inv_iter %>% g_k_iter)
    eps_hat_k <- epsilon_hat[i] + as.numeric((-1) * (X_train[i,] %>% delta_hat[[i]]))
    epsilon_hat_neg_k <- append(epsilon_hat_neg_k, eps_hat_k)
    loss_iter <- log(1 + exp(eps_hat_k)) + log(1 + exp(-1 * eps_hat_k))
    loss <- append(loss, loss_iter)
  }

  l_lambda <- append(l_lambda, mean(loss))
}

# Getting best lambda and then ERR_test

lambda_hat <- 10^iters[which.min(l_lambda)]

beta_hat_new <- minimize.robust.ridge(X_train, y_train, lambda = lambda_hat)

error <- y_test - (X_test %>% beta_hat_new)

ERR_test <- median(abs(error))

```

```

# Fitting the least squares ridge and getting the error differences

iters <- seq(log10(2), log10(50), length.out = 25)

lambdas <- 10^iters

ERR_test_ls_lambda <- c()

for (i in 1:length(lambdas)){
  XTX <- t(X_train) %*% X_train

  XTXlambda <- XTX + (lambdas[i] * diag(dim(X_train)[2]))

  XTXlambda_inv <- solve(XTXlambda)

  beta_hat_ls <- (XTXlambda_inv) %*% (t(X_train) %*% y_train)

  error_ls <- y_test - (X_test %*% beta_hat_ls)

  ERR_test_ls <- median(abs(error_ls))

  ERR_test_ls_lambda <- append(ERR_test_ls_lambda, ERR_test_ls)
}

ERR_diff <- ERR_test_ls_lambda - ERR_test

ERR_diff_0 <- ERR_diff

plot_data_3 <- as.data.frame(cbind(lambdas, ERR_diff_0))


plot_3 <- ggplot(plot_data_3,
  aes(x = lambdas, y = ERR_diff_0)) +
  geom_point(aes(x = lambdas, y = ERR_diff_0))+
  geom_line(aes(x = lambdas, y = ERR_diff_0))+
  xlab("log(lambda)") +
  ylab("ERR_test_ls - ERR_test") +

  theme_bw() +
  theme(axis.text.x=element_text(size=8),
    panel.grid.major = element_blank(),
    panel.grid.minor = element_blank())+
  ggtitle("Plot 3. Error difference against log(lambda), 0 outliers")
# geom_text(
#   label=round(plot_data$avg_gap,5),
#   vjust = -3.0,
#   hjust = 1.0,
#   position = position_dodge(width=1),
#   size =2
# )

```

plot_3

Plot 3. Error difference against log(lambda), 0 outliers



We see that with 0 outliers, there is no observable pattern since no robustness is needed. In fact, for some values of $\log(\lambda)$, the least squares ridge regression performs better than the tuned robust ridge regression which is not surprising, as no robustness is needed. It seems that for very high values of $\log(\lambda)$, the least squares ridge regression is underfitting.

5 outliers:

```
# Generating the data
set.seed(2022)
data <- generate.data(num.outliers = 5)

# Getting the training and test matrices
X_train <- as.matrix(data[[1]])
y_train <- as.matrix(data[[2]])
X_test  <- as.matrix(data[[3]])
y_test  <- as.matrix(data[[4]])

# Tuning the robust ridge regression
iters <- seq(-2, 1, length.out = 25)

lambdas <- 10^iters

l_lambda <- c()
```



```

for (j in 1:length(lambdas)){
  beta_hat <- minimize.robust.ridge(X_train, y_train, lambda = lambdas[j])

  epsilon_hat <- y_train - (X_train %%% beta_hat)

  l_prime <- ((exp(epsilon_hat)) / (1+exp(epsilon_hat))) - ((exp(-1 * epsilon_hat)) / ((1+exp(-1 * epsilon_hat)))
  l_2prime <- ((exp(epsilon_hat)) / ((1+exp(epsilon_hat))^2)) + ((exp(-1 * epsilon_hat)) / ((1+exp(-1 * epsilon_hat))^2))

  XTL <- (t(X_train)) %%% (diag(as.vector(l_2prime)))

  H <- ((1) / (dim(X_train)[1])) * (XTL %%% X_train) + (lambdas[j] * diag(dim(X_train)[2]))

  H_inv <- solve(H)

  # H_k <- list()
  H_k_inv <- list()
  g_k <- list()
  delta_hat <- list()
  epsilon_hat_neg_k <- c()
  loss <- c()

  for (i in 1:dim(X_train)[1]){
    numerator <- ((l_2prime[i]) / dim(X_train)[1]) * (H_inv %%% X_train[i,]) %%% (t(X_train[i,]) %%% H_inv)
    denominator <- as.numeric(1 - ((l_2prime[i]/dim(X_train)[1]) * (t(X_train[i,]) %%% H_inv) %%% (X_train[i,]) %%% t(X_train[i,])))
    # H_k_iter <- H - ( (1 / dim(X_train)[1]) * (l_2prime[i]) * (X_train[i,] %%% t(X_train[i,])) )
    # H_k[[i]] <- H_k_iter
    H_k_inv_iter <- H_inv + (numerator/denominator)
    H_k_inv[[i]] <- H_k_inv_iter
    g_k_iter <- (l_prime[i]/(dim(X_train)[1])) * (X_train[i,])
    g_k[[i]] <- g_k_iter
    delta_hat[[i]] <- (-1) * (H_k_inv_iter %%% g_k_iter)
    eps_hat_k <- epsilon_hat[i] + as.numeric((-1) * (X_train[i,] %%% delta_hat[[i]]))
    epsilon_hat_neg_k <- append(epsilon_hat_neg_k, eps_hat_k)
    loss_iter <- log(1 + exp(eps_hat_k)) + log(1 + exp(-1 * eps_hat_k))
    loss <- append(loss, loss_iter)
  }

  l_lambda <- append(l_lambda, mean(loss))
}

# Getting best lambda and then ERR_test

lambda_hat <- 10^iters[which.min(l_lambda)]

beta_hat_new <- minimize.robust.ridge(X_train, y_train, lambda = lambda_hat)

error <- y_test - (X_test %%% beta_hat_new)

ERR_test <- median(abs(error))

```

```

# Fitting the least squares ridge and getting the error differences

iters <- seq(log10(2), log10(50), length.out = 25)

lambdas <- 10^iters

ERR_test_ls_lambda <- c()

for (i in 1:length(lambdas)){
  XTX <- t(X_train) %*% X_train

  XTXlambda <- XTX + (lambdas[i] * diag(dim(X_train)[2]))

  XTXlambda_inv <- solve(XTXlambda)

  beta_hat_ls <- (XTXlambda_inv) %*% (t(X_train) %*% y_train)

  error_ls <- y_test - (X_test %*% beta_hat_ls)

  ERR_test_ls <- median(abs(error_ls))

  ERR_test_ls_lambda <- append(ERR_test_ls_lambda, ERR_test_ls)
}

ERR_diff <- ERR_test_ls_lambda - ERR_test

ERR_diff_5 <- ERR_diff

plot_data_4 <- as.data.frame(cbind(lambdas, ERR_diff_5))


plot_4 <- ggplot(plot_data_4,
  aes(x = lambdas, y = ERR_diff_5)) +
  geom_point(aes(x = lambdas, y = ERR_diff_5))+
  geom_line(aes(x = lambdas, y = ERR_diff_5))+
  xlab("log(lambda)") +
  ylab("ERR_test_ls - ERR_test") +

  theme_bw() +
  theme(axis.text.x=element_text(size=8),
    panel.grid.major = element_blank(),
    panel.grid.minor = element_blank())+
  ggtitle("Plot 4. Error difference against log(lambda), 5 outliers")
  # geom_text(
  #   label=round(plot_data$avg_gap,5),
  #   vjust = -3.0,
  #   # hjust = 1.0,
  #   position = position_dodge(width=1),
  #   size =2
  # )

```

plot_4

Plot 4. Error difference against $\log(\lambda)$, 5 outliers



In this situation, we see that the least squares ridge regression again performs better than the tuned robust ridge regression for some values of $\log(\lambda)$. Since there are only 5 outliers, perhaps the tuning is causing underfitting to some extent. Again for very large values of $\log(\lambda)$, the least squares ridge regression starts to underfit.

We already got the plot for 10 outliers in part d).

15 outliers:

```
# Generating the data
set.seed(2022)
data <- generate.data(num.outliers = 15)

# Getting the training and test matrices
X_train <- as.matrix(data[[1]])
y_train <- as.matrix(data[[2]])
X_test <- as.matrix(data[[3]])
y_test <- as.matrix(data[[4]])

# Tuning the robust ridge regression
iters <- seq(-2, 1, length.out = 25)

lambdas <- 10^iters

l_lambda <- c()
```

```

for (j in 1:length(lambdas)){
  beta_hat <- minimize.robust.ridge(X_train, y_train, lambda = lambdas[j])

  epsilon_hat <- y_train - (X_train %%% beta_hat)

  l_prime <- ((exp(epsilon_hat)) / (1+exp(epsilon_hat))) - ((exp(-1 * epsilon_hat)) / ((1+exp(-1 * epsilon_hat)))
  l_2prime <- ((exp(epsilon_hat)) / ((1+exp(epsilon_hat))^2)) + ((exp(-1 * epsilon_hat)) / ((1+exp(-1 * epsilon_hat))^2))

  XTL <- (t(X_train)) %%% (diag(as.vector(l_2prime)))

  H <- ((1) / (dim(X_train)[1])) * (XTL %%% X_train) + (lambdas[j] * diag(dim(X_train)[2]))

  H_inv <- solve(H)

  # H_k <- list()
  H_k_inv <- list()
  g_k <- list()
  delta_hat <- list()
  epsilon_hat_neg_k <- c()
  loss <- c()

  for (i in 1:dim(X_train)[1]){
    numerator <- ((l_2prime[i]) / dim(X_train)[1]) * (H_inv %%% X_train[i,]) %%% (t(X_train[i,]) %%% H_inv)
    denominator <- as.numeric(1 - ((l_2prime[i]/dim(X_train)[1]) * (t(X_train[i,]) %%% H_inv) %%% (X_train[i,]) %%% H_inv))
    # H_k_iter <- H - ( (1 / dim(X_train)[1]) * (l_2prime[i]) * (X_train[i,] %%% t(X_train[i,])) )
    # H_k[[i]] <- H_k_iter
    H_k_inv_iter <- H_inv + (numerator/denominator)
    H_k_inv[[i]] <- H_k_inv_iter
    g_k_iter <- (l_prime[i]/(dim(X_train)[1])) * (X_train[i,])
    g_k[[i]] <- g_k_iter
    delta_hat[[i]] <- (-1) * (H_k_inv_iter %%% g_k_iter)
    eps_hat_k <- epsilon_hat[i] + as.numeric((-1) * (X_train[i,] %%% delta_hat[[i]]))
    epsilon_hat_neg_k <- append(epsilon_hat_neg_k, eps_hat_k)
    loss_iter <- log(1 + exp(eps_hat_k)) + log(1 + exp(-1 * eps_hat_k))
    loss <- append(loss, loss_iter)
  }

  l_lambda <- append(l_lambda, mean(loss))
}

# Getting best lambda and then ERR_test

lambda_hat <- 10^iters[which.min(l_lambda)]

beta_hat_new <- minimize.robust.ridge(X_train, y_train, lambda = lambda_hat)

error <- y_test - (X_test %%% beta_hat_new)

ERR_test <- median(abs(error))

```

```

# Fitting the least squares ridge and getting the error differences

iters <- seq(log10(2), log10(50), length.out = 25)

lambdas <- 10^iters

ERR_test_ls_lambda <- c()

for (i in 1:length(lambdas)){
  XTX <- t(X_train) %*% X_train

  XTXlambda <- XTX + (lambdas[i] * diag(dim(X_train)[2]))

  XTXlambda_inv <- solve(XTXlambda)

  beta_hat_ls <- (XTXlambda_inv) %*% (t(X_train) %*% y_train)

  error_ls <- y_test - (X_test %*% beta_hat_ls)

  ERR_test_ls <- median(abs(error_ls))

  ERR_test_ls_lambda <- append(ERR_test_ls_lambda, ERR_test_ls)
}

ERR_diff <- ERR_test_ls_lambda - ERR_test

ERR_diff_15 <- ERR_diff

plot_data_5 <- as.data.frame(cbind(lambdas, ERR_diff_15))


plot_5 <- ggplot(plot_data_5,
  aes(x = lambdas, y = ERR_diff_15)) +
  geom_point(aes(x = lambdas, y = ERR_diff_15))+
  geom_line(aes(x = lambdas, y = ERR_diff_15))+
  xlab("log(lambda)") +
  ylab("ERR_test_ls - ERR_test") +

  theme_bw() +
  theme(axis.text.x=element_text(size=8),
    panel.grid.major = element_blank(),
    panel.grid.minor = element_blank())+
  ggtitle("Plot 5. Error difference against log(lambda), 15 outliers")
  # geom_text(
  #   label=round(plot_data$avg_gap,5),
  #   vjust = -3.0,
  #   # hjust = 1.0,
  #   position = position_dodge(width=1),
  #   size =2
  # )

```

plot_5

Plot 5. Error difference against log(lambda), 15 outliers



Again, for very small values of $\log(\lambda)$, the least squares ridge regression seems to be performing better, and then starts to underfit. However, as $\log(\lambda)$ increases beyond approximately 25, we see that the least squares ridge regression starts to get closer to the tuned robust ridge regression in terms of median prediction error.

20 outliers:

```
# Generating the data
set.seed(2022)
data <- generate.data(num.outliers = 20)

# Getting the training and test matrices
X_train <- as.matrix(data[[1]])
y_train <- as.matrix(data[[2]])
X_test  <- as.matrix(data[[3]])
y_test  <- as.matrix(data[[4]])

# Tuning the robust ridge regression
iters <- seq(-2, 1, length.out = 25)

lambdas <- 10^iters

l_lambda <- c()

for (j in 1:length(lambdas)){
```

```

beta_hat <- minimize.robust.ridge(X_train, y_train, lambda = lambdas[j])

epsilon_hat <- y_train - (X_train %>% beta_hat)

l_prime <- ((exp(epsilon_hat)) / (1+exp(epsilon_hat))) - ((exp(-1 * epsilon_hat)) / ((1+exp(-1 * epsilon_hat)))
l_2prime <- ((exp(epsilon_hat)) / ((1+exp(epsilon_hat))^2)) + ((exp(-1 * epsilon_hat)) / ((1+exp(-1 * epsilon_hat))^2))

XTL <- (t(X_train)) %>% (diag(as.vector(l_2prime)))

H <- ((1) / (dim(X_train)[1])) * (XTL %>% X_train) + (lambdas[j] * diag(dim(X_train)[2]))

H_inv <- solve(H)

# H_k <- list()
H_k_inv <- list()
g_k <- list()
delta_hat <- list()
epsilon_hat_neg_k <- c()
loss <- c()

for (i in 1:dim(X_train)[1]){
  numerator <- ((l_2prime[i] / dim(X_train)[1]) * (H_inv %>% X_train[i,]) %>% (t(X_train[i,]) %>% H_inv))
  denominator <- as.numeric(1 - ((l_2prime[i]/dim(X_train)[1]) * (t(X_train[i,]) %>% H_inv) %>% (X_train[i,] %>% t(X_train[i,]))))
  # H_k_iter <- H - ( (1 / dim(X_train)[1]) * (l_2prime[i]) * (X_train[i,] %>% t(X_train[i,])) )
  # H_k[[i]] <- H_k_iter
  H_k_inv_iter <- H_inv + (numerator/denominator)
  H_k_inv[[i]] <- H_k_inv_iter
  g_k_iter <- (l_prime[i]/(dim(X_train)[1])) * (X_train[i,])
  g_k[[i]] <- g_k_iter
  delta_hat[[i]] <- (-1) * (H_k_inv_iter %>% g_k_iter)
  eps_hat_k <- epsilon_hat[i] + as.numeric((-1) * (X_train[i,] %>% delta_hat[[i]]))
  epsilon_hat_neg_k <- append(epsilon_hat_neg_k, eps_hat_k)
  loss_iter <- log(1 + exp(eps_hat_k)) + log(1 + exp(-1 * eps_hat_k))
  loss <- append(loss, loss_iter)
}

l_lambda <- append(l_lambda, mean(loss))
}

# Getting best lambda and then ERR_test

lambda_hat <- 10^iters[which.min(l_lambda)]

beta_hat_new <- minimize.robust.ridge(X_train, y_train, lambda = lambda_hat)

error <- y_test - (X_test %>% beta_hat_new)

ERR_test <- median(abs(error))

# Fitting the least squares ridge and getting the error differences

```

```

iters <- seq(log10(2), log10(50), length.out = 25)

lambdas <- 10^iters

ERR_test_ls_lambda <- c()

for (i in 1:length(lambdas)){
  XTX <- t(X_train) %*% X_train

  XTXlambda <- XTX + (lambdas[i] * diag(dim(X_train)[2]))

  XTXlambda_inv <- solve(XTXlambda)

  beta_hat_ls <- (XTXlambda_inv) %*% (t(X_train) %*% y_train)

  error_ls <- y_test - (X_test %*% beta_hat_ls)

  ERR_test_ls <- median(abs(error_ls))

  ERR_test_ls_lambda <- append(ERR_test_ls_lambda, ERR_test_ls)
}

ERR_diff <- ERR_test_ls_lambda - ERR_test

ERR_diff_20 <- ERR_diff

plot_data_6 <- as.data.frame(cbind(lambdas, ERR_diff_20))


plot_6 <- ggplot(plot_data_6,
  aes(x = lambdas, y = ERR_diff_20)) +
  geom_point(aes(x = lambdas, y = ERR_diff_20))+
  geom_line(aes(x = lambdas, y = ERR_diff_20))+
  xlab("log(lambda)") +
  ylab("ERR_test_ls - ERR_test") +

  theme_bw() +
  theme(axis.text.x=element_text(size=8),
    panel.grid.major = element_blank(),
    panel.grid.minor = element_blank())+
  ggtitle("Plot 6. Error difference against log(lambda), 20 outliers")
  # geom_text(
  #   label=round(plot_data$avg_gap,5),
  #   vjust = -3.0,
  #   # hjust = 1.0,
  #   position = position_dodge(width=1),
  #   size =2
  # )

plot_6

```


Plot 6. Error difference against $\log(\lambda)$, 20 outliers



For 20 outliers, we see a clear pattern. As the regularization for the least squares ridge regression increases, the median absolute prediction error approaches that of the tuned robust ridge regression, and for small values of $\log(\lambda)$, the error difference is much higher than in plots seen for fewer outliers.

Finally, 25 outliers:

```
# Generating the data
set.seed(2022)
data <- generate.data(num.outliers = 25)

# Getting the training and test matrices
X_train <- as.matrix(data[[1]])
y_train <- as.matrix(data[[2]])
X_test <- as.matrix(data[[3]])
y_test <- as.matrix(data[[4]])

# Tuning the robust ridge regression
iters <- seq(-2, 1, length.out = 25)

lambdas <- 10^iters

l_lambda <- c()

for (j in 1:length(lambdas)){
  beta_hat <- minimize.robust.ridge(X_train, y_train, lambda = lambdas[j])

  epsilon_hat <- y_train - (X_train %*% beta_hat)
```

```

l_prime <- ((exp(epsilon_hat)) / (1+exp(epsilon_hat))) - ((exp(-1 * epsilon_hat)) / ((1+exp(-1 * epsilon_hat)))
l_2prime <- ((exp(epsilon_hat)) / ((1+exp(epsilon_hat))^2)) + ((exp(-1 * epsilon_hat)) / ((1+exp(-1 * epsilon_hat)))

XTL <- (t(X_train)) %*% (diag(as.vector(l_2prime)))

H <- ((1) / (dim(X_train)[1])) * (XTL %*% X_train) + (lambdas[j] * diag(dim(X_train)[2]))

H_inv <- solve(H)

# H_k <- list()
H_k_inv <- list()
g_k <- list()
delta_hat <- list()
epsilon_hat_neg_k <- c()
loss <- c()

for (i in 1:dim(X_train)[1]){
  numerator <- ((l_2prime[i]) / dim(X_train)[1]) * (H_inv %*% X_train[i,]) %*% (t(X_train[i,]) %*% H_inv)
  denominator <- as.numeric(1 - ((l_2prime[i]/dim(X_train)[1]) * (t(X_train[i,]) %*% H_inv) %*% (X_train[i,])))
  # H_k_iter <- H - ( (1 / dim(X_train)[1]) * (l_2prime[i]) * (X_train[i,] %*% t(X_train[i,]))) )
  # H_k[[i]] <- H_k_iter
  H_k_inv_iter <- H_inv + (numerator/denominator)
  H_k_inv[[i]] <- H_k_inv_iter
  g_k_iter <- (l_prime[i]/(dim(X_train)[1])) * (X_train[i,])
  g_k[[i]] <- g_k_iter
  delta_hat[[i]] <- (-1) * (H_k_inv_iter %*% g_k_iter)
  eps_hat_k <- epsilon_hat[i] + as.numeric((-1) * (X_train[i,] %*% delta_hat[[i]]))
  epsilon_hat_neg_k <- append(epsilon_hat_neg_k, eps_hat_k)
  loss_iter <- log(1 + exp(eps_hat_k)) + log(1 + exp(-1 * eps_hat_k))
  loss <- append(loss, loss_iter)
}

l_lambda <- append(l_lambda, mean(loss))
}

# Getting best lambda and then ERR_test

lambda_hat <- 10^iters[which.min(l_lambda)]

beta_hat_new <- minimize.robust.ridge(X_train, y_train, lambda = lambda_hat)

error <- y_test - (X_test %*% beta_hat_new)

ERR_test <- median(abs(error))

# Fitting the least squares ridge and getting the error differences

iters <- seq(log10(2), log10(50), length.out = 25)

lambdas <- 10^iters

```

```

ERR_test_ls_lambda <- c()

for (i in 1:length(lambdas)){
  XTX <- t(X_train) %*% X_train

  XTXlambda <- XTX + (lambdas[i] * diag(dim(X_train)[2]))

  XTXlambda_inv <- solve(XTXlambda)

  beta_hat_ls <- (XTXlambda_inv) %*% (t(X_train) %*% y_train)

  error_ls <- y_test - (X_test %*% beta_hat_ls)

  ERR_test_ls <- median(abs(error_ls))

  ERR_test_ls_lambda <- append(ERR_test_ls_lambda, ERR_test_ls)
}

ERR_diff <- ERR_test_ls_lambda - ERR_test

ERR_diff_25 <- ERR_diff

plot_data_7 <- as.data.frame(cbind(lambdas, ERR_diff_25))


plot_7 <- ggplot(plot_data_7,
  aes(x = lambdas, y = ERR_diff_25)) +
  geom_point(aes(x = lambdas, y = ERR_diff_25))+
  geom_line(aes(x = lambdas, y = ERR_diff_25))+
  xlab("log(lambda)") +
  ylab("ERR_test_ls - ERR_test") +

  theme_bw() +
  theme(axis.text.x=element_text(size=8),
    panel.grid.major = element_blank(),
    panel.grid.minor = element_blank())+
  ggtitle("Plot 7. Error difference against log(lambda), 25 outliers")
  # geom_text(
  #   label=round(plot_data$avg_gap,5),
  #   vjust = -3.0,
  #   hjust = 1.0,
  #   position = position_dodge(width=1),
  #   size =2
  # )

plot_7

```

Plot 7. Error difference against $\log(\lambda)$, 25 outliers



The plot obtained for 25 outliers is similar to the one for 20 outliers, as beyond a some deviations at the beginning, as the regularization for least squares increases, again it approaches the tuned robust ridge regression in terms of median absolute prediction error.